

# HIVE

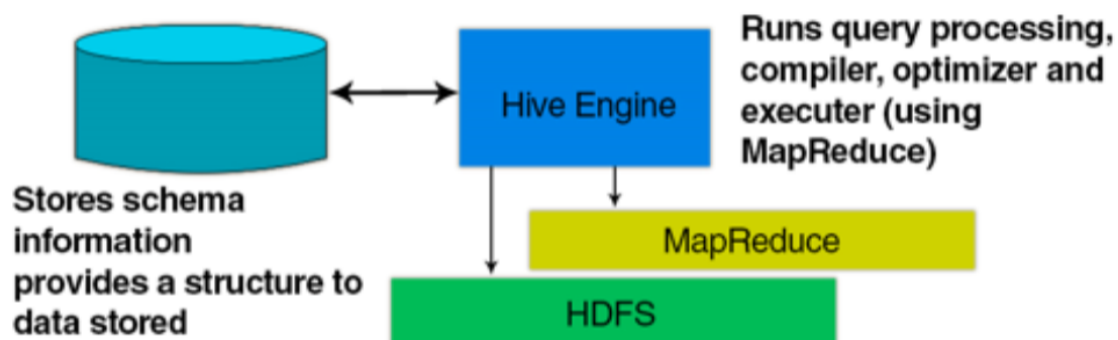
## What is HIVE ?

- HIVE is a query interface on top of Hadoop's native Map-Reduce
- HIVE is a data warehouse
- HIVE allows users to write SQL style queries in a native language known as Hive Query Language (HQL)
- HIVE execution engine converts the scripts written in HQL into JAR files (map reduce) to execute in the cluster
- HIVE reads data from HDFS
- Allows creation of tables to operate on structured data
- The table's schema information (table metadata) is saved in HIVE metastore which is borrowed from an RDBMS (Derby is default database)
- HIVE is not an RDBMS

## Why HIVE ?


- Hadoop is known for its Map-Reduce engine for parallelizing data processing operations using HDFS as its native file storage system
- Map Reduce does not provide user friendly libraries or interfaces to deal with unstructured data handling
- Very tight dependency of JAVA if one needs to use the Map-Reduce framework
- An operation like left inner join would need around 200-300 lines of code in JAVA Map-Reduce whereas in SQL it would just be a couple of lines of code
- Analysts from SQL experience of having come from RDBMS world and DW/BI world cannot program in JAVA in order to use
- To enable SQL developers to exploit the power of Hadoop, an abstraction interface was developed on top of native Map-Reduce
- This interface (engine) was called HIVE and was officially developed by Facebook and initial release was in the year 2010

## Architectural Overview

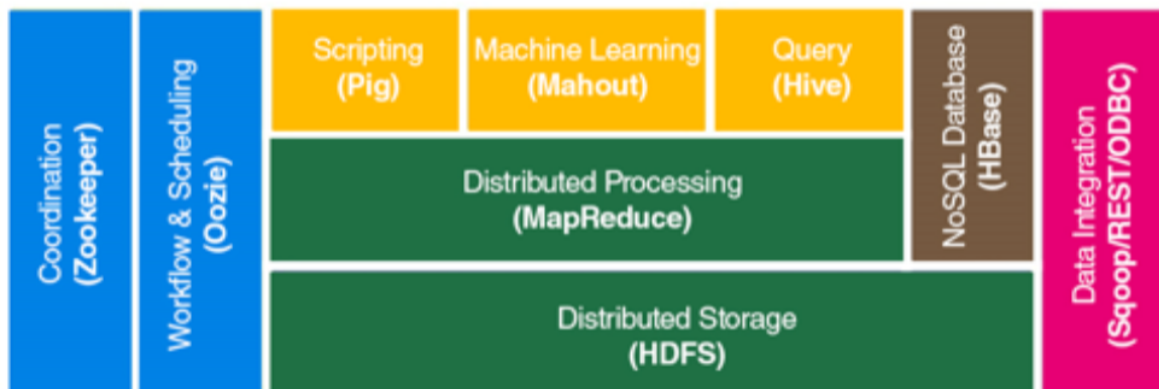


## Working of HIVE

- Hive allows a way to project a table structure on the data in HDFS (structured data in HDFS)
- The table metadata is saved separately from the data
- In reality, we do not actually load the data into the place where HIVE tables are created
- HIVE table information (table metadata is saved in meta store)



How does HIVE fit into the Hadoop ecosystem?



Things HIVE cannot do efficiently

- Ad hoc real time queries
- OLTP (Online Line Transaction Processing)
- No ACID support (ACID support is limited)
- Not suited for frequent updates and inserts (inserts and updates are allowed in recent releases of HIVE)
- Not recommended for small data sets
- Not meant for unstructured data analysis

# Installation

References :

[https://www.tutorialspoint.com/hive/hive\\_installation.htm](https://www.tutorialspoint.com/hive/hive_installation.htm)

<https://drive.google.com/drive/folders/1XdPbyAc9iWml0fPPNX91Yq3BRwkZAG2M>

```
sudo apt update
sudo apt install openjdk-8-jdk -y
```

```
sudo nano ~/.bashrc
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
export PATH=$PATH:$JAVA_HOME/bin
source ~/.bashrc
```

```
java -version; javac -version
```

```
sudo apt install openssh-server openssh-client -y
ssh-keygen -t rsa -P "" -f ~/.ssh/id_rsa
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
chmod 0600 ~/.ssh/authorized_keys
ssh localhost
(sudo service ssh start)
```

# Hadoop -

<https://d1cdn.apache.org/hadoop/common/hadoop-3.2.3/hadoop-3.2.3.tar.gz>

```
tar xzf hadoop-3.2.3.tar.gz
# put it in /usr/local/ as hadoop
```

# Editing 6 important files

```
sudo nano ~/.bashrc
export HADOOP_HOME=/usr/local/hadoop
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export YARN_HOME=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export PATH=$PATH:$HADOOP_HOME/sbin:$HADOOP_HOME/bin
export HADOOP_OPTS="-Djava.library.path=$HADOOP_HOME/lib/native"
export HADOOP_OPTS="$HADOOP_OPTS -Djava.library.path=$HADOOP_HOME/lib/native"
source ~/.bashrc
```

```
cd $HADOOP_HOME/etc/hadoop
```

```
sudo nano core-site.xml
<configuration>
  <property>
    <name>fs.default.name</name>
```

```
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

sudo nano hdfs-site.xml

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
  <property>
    <name>dfs.name.dir</name>
    <value>file:///home/archit/hadoopinfra/hdfs/namenode </value>
  </property>
  <property>
    <name>dfs.data.dir</name>
    <value>file:///home/archit/hadoopinfra/hdfs/datanode </value>
  </property>
</configuration>
```

sudo nano yarn-site.xml

```
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
</configuration>
```

sudo nano mapred-site.xml

```
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>
```

# Verifying Hadoop Installation

cd ~

hdfs namenode -format

start-dfs.sh

start-yarn.sh

<http://localhost:9870/>

# HIVE -

<https://dlcdn.apache.org/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz>

```
tar xzf apache-hive-3.1.2-bin.tar.gz
# put it in /usr/local/ as hive
```

```
sudo nano ~/.bashrc
```

```
export HIVE_HOME=/usr/local/hive
export PATH=$PATH:$HIVE_HOME/bin
export CLASSPATH=$CLASSPATH:/usr/local/Hadoop/lib/*:.
export CLASSPATH=$CLASSPATH:/usr/local/hive/lib/*:.
source ~/.bashrc
```

```
sudo nano $HIVE_HOME/bin/hive-config.sh
export HADOOP_HOME=/usr/local/hadoop
```

```
# Create Hive directories in HDFS
hdfs dfs -mkdir /tmp
hdfs dfs -chmod g+w /tmp
hdfs dfs -mkdir -p /user/hive/warehouse
hdfs dfs -chmod g+w /user/hive/warehouse
```

```
# Fixing guava problem
rm $HIVE_HOME/lib/guava-19.0.jar
cp /usr/local/hadoop/share/hadoop/hdfs/lib/guava-27.0-jre.jar $HIVE_HOME/lib/
```

```
# Initialize Derby and hive
schematool -initSchema -dbType derby
```

```
hive
```

```
optional Step – Edit hive-site.xml
=====
cd $HIVE_HOME/conf
cp hive-default.xml.template hive-
site.xml
```

```
sudo nano hive-site.xml – change
metastore location to above created hdfs
path(/user/hive/warehouse)
```

```
# Starting HIVE:
sudo service ssh start
start-dfs.sh
start-yarn.sh
Hive
```

## Basics

# List contents of HDFS home directory

```
hadoop fs -ls /
```

# Make directory

```
hadoop fs -mkdir foo
```

```
cd /home/archit/myhivedata
```

# Launch in the directory that we want to use

```
hive
```

# Create database

```
create database vin_emp;
```

```
use vin_emp;
```

# Create table

```
create table emp_global(id int,name string,city string,continent string)
```

```
row format delimited          # Data present row after row
```

```
fields terminated by ','      # Column elements separated by ,
```

```
stored as textfile;           # How HIVE internally stores the file(txt-default)
```

```
show tables;
```

# Load data into table

# (the txt file must be in the dir from where we launched HIVE or absolute path)

```
load data local inpath 'employees.txt' into table emp_global;
```

^ to specify data present in local FS(not HDFS, skip local)

```
select * from emp_global;
```

```
describe emp_global;          # show schema of table
```

```
describe extended emp_global;
```

```
# TableType:MANAGED_TABLE     # or internal tables
```

# (HIVE engine will take total control of data in HDFS, drop table; data in HDFS lost)

```
# TableType:EXTERNAL_TABLE
```

# Where is the database created, where the data is stored?

# def. location where HIVE engine is configured to store/create databases & tables

```
hadoop fs -ls /user/hive/warehouse
```

# this contains vin\_emp.db (.db auto generated to show it is a database folder)

```
hadoop fs -ls /user/hive/warehouse/vin_emp.db
```

```
# this contains emp_global (table directory containing table data)

# Delete table
drop table emp_global;
# Delete database
drop database vin_emp;          # error : vin_emp is not empty
drop database vin_emp cascade;  # drop all tables the drop database

# Create database other than default warehouse
create database vin_emp_loc location '/user/archit/myhivedata';
                                ^- should exist in HDFS

show databases;
```

## External tables in hive

External tables point to the location in HDFS where the data is already present unlike loading the data like in managed/internal tables.

# copy file from local FS into HDFS location

hadoop fs -mkdir empdata # creating a folder in HDFS home directory

hadoop fs -ls # check

hadoop fs -put empglobal.csv empdata

# empglobal.csv should be in the pwd & empdata is the folder

hadoop fs -cat empdata/empglobal.csv # check file

# create an external table

create database emp\_ext location '/user/archit/myhive\_ext';

hadoop fs -ls # check created directory

create external table if not exists emp\_global(id int,name string,city string,continent string)

^- now ^- optional

row format delimited

fields terminated by ','

stored as textfile

location '/user/archit/empdata';

^- containing empglobal.csv

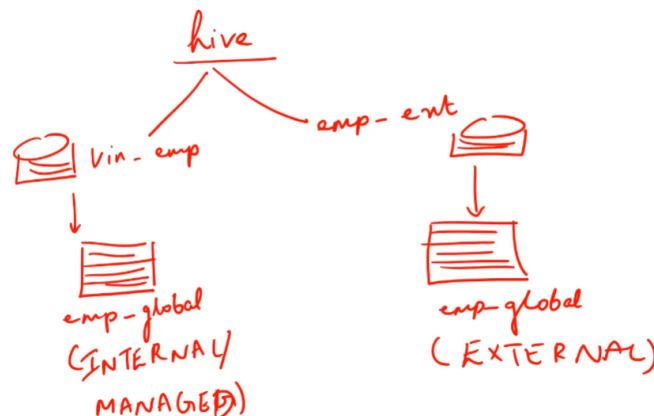
(spec. that table(engine) must look in that HDFS folder for data)

**NOTE:** We do NOT load data in the table, instead we create tables with the same structure in which data is already present in the HDFS location and make the table point to that location. Then we can issue queries directly without loading.

select \* from emp\_global;

describe extended emp\_global;

# TableType:EXTERNAL\_TABLE





```
# Internal table
use vin_emp;
drop table emp_global;
hadoop fs -ls /user/archit/warehouse/vin_emp.db      # Empty
# corresponding HDFS gets auto deleted, HIVE takes total control of data as well
```

```
# External tables
use emp_ext;
drop table emp_global;
hadoop fs -ls empdata
# still have file empglobal.csv
```

Advantages -

- 1) Accidental dropping of data in HIVE will not result in the deletion of actual data present in HDFS.
- 2) If data belongs to some other team, we can still use data to perform queries by just creating table structure which point to that data.

## Loading different file formats

HIVE supported file formats -

- RC (Row Column) FF
- ORC (Optimized Row Column) FF
- Sequence FF
- Avro FF
- Parquet FF

```
create table emp_global(id int,name string,city string,continent string)
row format delimited          # Data present row after row
fields terminated by ','      # Column elements separated by ,
stored as textfile;          # How HIVE internally stores the file(txt-default)
```

### # ORC

```
create table emp_global_orc(id int,name string,city string,continent string)
row format delimited
fields terminated by ','
stored as orcfile;
```

# Loading data into ORC table done indirectly, we run select \* query (fetch data) from an existing table which is not in ORC FF (actually saved in text FF), then we try to push the result of the query into the ORC table.

```
insert into table emp_global_orc select * from emp_global;
```

```
select * from emp_global_orc;    # same data
```

# internal storage and file format varies

```
hadoop fs -ls /user/hive/warehouse/vin_emp.db/emp_global
# contains employees.txt (109 bytes)
```

```
hadoop fs -ls /user/hive/warehouse/vin_emp.db/emp_global_orc
# contains 000000_0 (644 bytes)
```

# ORC format supports ACID operations

```
hadoop fs -cat /user/hive/warehouse/vin_emp.db/emp_global_orc/000000_0
# cannot directly makes sense of what is present (looks encrypted)
```

### # Sequence FF

```
create table emp_global_seq LIKE emp_global_orc stored as sequencefile;
                        ^- to clone table schema           ^- saved as
```

```
insert into table emp_global_seq select * from emp_global;
```

```
hadoop fs -ls /user/hive/warehouse/vin_emp.db/emp_global_seq  
# contains 000000_0 (251 bytes)
```

# Avro FE

```
create table emp_global_avro LIKE emp_global_orc stored as avrofile;  
insert into table emp_global_avro select * from emp_global;
```

# Parquet FE

```
create table emp_global_par LIKE emp_global_orc stored as parquetfile;  
insert into table emp_global_par select * from emp_global;
```

## Loading data into Hive tables

Load data into a table in HIVE, especially when data is not just in the form of rows and columns. (Complex data structure, multiple fields[array in 1 single column], structure within a column element)

NOTE: schema check in HIVE does take place while writing, it checks while querying. Operations on HIVE are called schema-on-read.

# Sample dataset (saved as /home/archit/myhivedata/siblings.txt)

Ram,43, India,Laxman

Alice,55,USA,Alex#Tim

Somu,22,India,Sham#Ram

Ted,67,Australia,Kevin#Jincy#Tanya#Kate

```
create table if not exists sibling_data (  
name string,  
age int,  
country string,  
siblings array<string>)  
row format delimited  
fields terminated by ','  
collection items terminated by '#'  
lines terminated by '\n'  
stored as textfile;
```

```
load data local inpath '/home/archit/myhivedata/siblings.txt' into table sibling_data;
```

# Structure item loading into HIVE

# Sample dataset (stored as /home/archit/myhivedata/autodata.txt)

Maruthi-Suzuki,Swift,Petrol,hatchback#5#5,1197#85.00

Maruthi-Suzuki,Swift,Diesel,hatchback#5#5,1248#83.14

Ranault,Duster,Diesel,compact-suv#5#6,1498#106

```
create table auto_details(company string, model string, fuel string,  
basic_specs struct<vehicle_type:string, doors:int, gears:int>,  
engine_specs struct<cc:int, bhp:double>)  
row format delimited  
fields terminated by ','  
collection items terminated by '#' ;
```

```
load data local inpath '/home/archit/myhivedata/autodata.txt' into table auto_details ;
```

## Simple Operations on Hive tables

Alter, Add, Drop columns; rename table

# Rename existing table

```
alter table auto_details rename to auto_table;
```

# hadoop fs -ls /user/hive/warehouse/vin\_emp.db/ (check name change)

# Rename column in table

```
alter table auto_details change fuel fuel_type string;
```

^old ^new ^data type(recomm. to retain og)

# Add new column

```
alter table auto_details add columns (mileage double);
```

^ new ^ data type

```
select * from auto_details; # NULL in mileage column
```

# Drop existing columns

```
alter table auto_details replace columns(company string, model string, fuel_type string);
```

^ name of all columns that need to be retained

## Query Operations on Hive tables

# Sample data at (/home/archit/myhivedata/empdata.txt)

emp\_id,emp\_name,salary,designation,department,age

564,Ram,300000,VP,Board,47

786,Bob,46000,Developer,Finance,32

987,Ravi,76000,Tech Lead,IT,43

001,Shiva,NA,CEO,Board,49

348,Aditya,100000,Project Manager,IT,46

567,Surya,6579,Intern,HR, 22

789,Sam,9877,Intern,IT,25

123,Asha,34000,HR,HR,32

121,Deepthi,45000,HR manager,HR,43

111,Krishna,40000,Manager,Infra,35

710,Soma,290600,QA,IT,26

698,Piyush,45000,QA Lead,IT,34

498,Manish,67000,Manager,Finance,48

302,Sita,200000,AVP,Board,42

002,Parvathy,NA,C00,Board,45

create database company;

create table if not exists company.empdata (

empid int, ^-specifying DB without 'use' command

empname string,

salary double,

designation string,

department string,

age int)

row format delimited

fields terminated by ','

lines terminated by '\n'

tblproperties('skip.header.line.count'='1');

^-builtin HIVE func ^-specifying header column no.'s

(large tables can have headers in multiple[4/5] lines)

# Incase if the file has a header and footer which comprises of first 2 and last 2 lines of the file, then use this syntax to ignore them

tblproperties('skip.header.line.count'='2', 'skip.footer.line.count'='2');

load data local inpath '/home/archit/myhivedata/empdata.txt' into table empdata;

null

## # Fundamental Queries -

# pick the rows of employees who work in HR department

```
select * from empdata where department="HR";
```

# pick rows of employees who work in the HR dept. & have salary more than 25000

```
select * from empdata where department="HR" and salary > 25000;
```

# select name and age of people in the HR dept. & have salary more than 25000

```
select empname, age from empdata where department="HR" and salary > 25000;
```

# pick employees in the increasing order of their salary

```
select * from empdata order by salary desc;
```

# count number of rows in the table

```
select count(*) from empdata;
```

# count number of employees in each department

```
select department, count(*) from empdata group by department;
```

# ignore columns that have NULL

```
select * from empdata where salary is not null;
```

# substring matching; pick employees whose designation is Manager/Lead

```
select * from empdata where designation rlike "Manager" OR designation rlike  
"Lead" OR designation rlike "manager";
```

# average salary grouped by department

```
select department, avg(salary) from empdata group by department;
```

## Querying complex structures

```
# query subfield of a column in case of complex structure (using auto_details)
select * from auto_details where engine_specs.bhp > 85;
```

### # Joins

```
# Sample table at (/home/archit/myhivedata/employee_epf.txt)
```

```
emp_id,emp_name,EPF
```

```
001,Shiva,False
```

```
002,Parvathy,False
```

```
348,Aditya,False
```

```
710,Soma,False
```

```
498,Manish,False
```

```
create table emp_epf
```

```
(empid int,empname string,epf boolean)
```

```
row format delimited
```

```
fields terminated by ','
```

```
tblproperties('skip.header.line.count'='1');
```

```
load data local inpath '/home/archit/myhivedata/employee_epf.txt' into table emp_epf;
```

```
SET hive.auto.convert.join=false;      # To enable joins
```

```
select emp.empname, emp.salary from emp_epf pf join empdata emp on (pf.empid = emp.empid);
```

```
          _____      _____      ^__ alias __^
          (emp_epf as pf)  (empdata as emp)
```

# No. of mappers = 2 and reducers = 1 when we run it. Since both data is smaller than block size in HDFS, if the table is spread across multiple blocks of data then mappers > 2.

# Default join in HIVE is inner join

# Replace 'join' by the following

left outer join

right outer join

full outer join



## Views

Pseudo tables that contain subset of a database or query which we do not want to re-run each time. They are logical tables which get automatically deleted when the parent table gets deleted.

# Creating a view

Create view if not exists high\_sal as select \* from empdata where salary > 50000;

show views;                      OR                      show tables;

describe formatted high\_sal;

# Table Type: VIRTUAL\_VIEW

drop table empdata;                      # Deleting parent table

show tables;                      # Still shows high\_sal

select \* from high\_sal;

# Table not found 'empdata' in definition of VIEW high\_sal;

# Segregating managerial and non-managerial employees into 2 views

# query to create a view containing the data of employees who are not managers, leads, CEO, VP, AVP and COO

create view regular\_employees as

select empid, empname, salary, department from empdata where designation not rlike "Manager" and designation not rlike "Lead" and designation not rlike "manager" and designation not rlike "VP" and designation not rlike "AVP" and designation not rlike "CEO" and designation not rlike "COO";

# query to create a view containing the data of employees who are managers, leads, CEO, VP, AVP and COO

create view mgmt\_employees as

select empid, empname, salary, department from empdata where designation rlike "Manager" or designation rlike "Lead" or designation rlike "manager" or designation rlike "VP" or designation rlike "AVP" or designation rlike "CEO" or designation not rlike "COO";

# we can now run any query on these virtual views treating them like tables

# Partitions

Dividing data based on a particular column

# Ex: Create table organized based on continent

```
create table emp_global_part(id int,name string,city string, country string)
partitioned by (continent string)
row format delimited
fields terminated by ','
stored as textfile;
```

# Load data into partitioned table

```
insert into table emp_global_part partition(continent="Asia") select * from emp_global
where continent = "Asia";
```

```
insert into table emp_global_part partition(continent="North America") select * from
emp_global where continent = "North America";
```

```
insert into table emp_global_part partition(continent="European Union") select * from
emp_global where continent = "European Union";
```

# PROJECT 1 : Yellow taxi trip analysis using Hive

archit.wadehra@gmail.com

## Creating table

```
create table taxi(vendor_id string, pickup_datetime string, dropoff_datetime
string, passenger_count int, trip_distance DECIMAL(9,6), pickup_longitude
DECIMAL(9,6), pickup_latitude DECIMAL(9,6), rate_code int, store_and_fwd_flag
string, dropoff_longitude DECIMAL(9,6), dropoff_latitude DECIMAL(9,6),
payment_type string, fare_amount DECIMAL(9,6), extra DECIMAL(9,6), mta_tax
DECIMAL(9,6), tip_amount DECIMAL(9,6), tolls_amount DECIMAL(9,6), total_amount
DECIMAL(9,6), trip_time_in_secs int)
row format delimited
fields terminated by ','
tblproperties('skip.header.line.count'='1');
```

## Loading data

```
load data local inpath '/home/archit/myhivedata/yellow_tripdata_2015 -01-06.csv'
into table taxi;
```

1) What is the total number of trips ( equal to the number of rows)?

```
select count(*) from taxi;
```

10000

2) What is the total revenue generated by all the trips? The fare is stored in the column total\_amount.

```
select sum(total_amount) from taxi;
```

160546.810000

3) What fraction of the total is paid for tolls? The toll is stored in tolls\_amount.

```
select sum(tolls_amount) / sum(total_amount) from taxi;
```

0.0155530340341237549

4) What fraction of it is driver tips? The tip is stored in tip\_amount.

```
select sum(tip_amount) / sum(total_amount) from taxi;
```

0.1078520339332808917

5) What is the average trip amount?

NOTE : No column is called "trip\_amount" assuming it to be "tip\_amount"

```
select AVG(tip_amount) from taxi;
```

1.731530000000000000000000000000

- 6) What is the average distance of the trips? Distance is stored in the column trip\_distance.

```
select AVG(trip_distance) from taxi;
```

3.253033000000000000000000000000

- 7) How many different payment types are used?

```
select count(distinct payment_type) from taxi;
```

4

- 8) For each payment type, display the following details:

- Average fare generated (fare\_amount)
- Average tip (tip\_amount)
- Average tax (mta\_tax)

```
select payment_type, avg(fare_amount), avg(tip_amount),  
avg(mta_tax) from taxi group by payment_type;
```

PT	Average fare amount	Average tip amount	Average tax amount
1	13.5610182726846790566921755	2.7042480087459003592066219	0.4971107293456192409807903
2	11.3933830985915492957746479	0.000000000000000000000000000000	0.4988732394366197183098592
3	13.2107894736842105263157895	0.000000000000000000000000000000	0.4210526315789473684210526
4	12.222222222222222222222222222222	0.000000000000000000000000000000	0.500000000000000000000000000000

- 9) On average which hour of the day generates the highest revenue?

```
select substr(pickup_datetime,12,2), avg(total_amount) from taxi  
group by substr(pickup_datetime,12,2);
```

Hour	Revenue
00	15.3199724061810154525386313
22	16.2401006214856466410180527
23	16.1093846153846153846153846

22 (22nd hour generates the highest average revenue)